

HPC Experiment 7 Report

Block Based Matrix Multiplication

Name: Reuben Skariah Mathew

Roll No: CED18I042

Programming Environment: OpenMP

Problem: Block Based Matrix Multiplication

Date: 2nd September 2021

Hardware Configuration:

Processor: Intel(R) Core(TM) i5-8250U CPU @ 1.60GHz

Sockets: 1

Cores per Socket: 4

Threads per Core: 2

L1 Cache: 32 kB

L2 Cache: 256 kB

L3 Cache: 6 MB

RAM: 8 GB

Serial Code:

```
#include <omp.h>
#include <stdio.h>
#include <stdlib.h>

#define MIN(X, Y) (((X) < (Y)) ? (X) : (Y))

int block[] = {5, 10, 50, 60, 70, 80, 90, 100};
int main()
{
    double st, et, tt[8];

    int L = 500, M = 500, N = 500;

    double matrix_x[500][500], matrix_y[500][500], matrix_z[500][500];

    for (int i = 0; i < 500; i++)
    {
        for (int j = 0; j < 500; j++)
        {
            matrix_x[i][j] = (i * 5.123);
            matrix_y[i][j] = (j * 3.123);
            matrix_z[i][j] = 0;
        }
    }

    for(int bl = 0; bl < 4; bl++)
    {
        st = omp_get_wtime();

        #pragma omp for
```

```

for(int jj = 0; jj < N; jj = jj + block[bl])
{
    for(int kk = 0; kk < N; kk = kk + block[bl])
    {
        for(int i = 0; i < N; i++)
        {
            for(int j = jj; j < MIN(jj + block[bl], N); j++)
            {
                double r = 0;
                for(int k = kk; k < MIN(kk + block[bl], N); k++)
                {
                    r += matrix_x[i][k] * matrix_y[k][j];
                }
                matrix_z[i][j] = matrix_z[i][j] + r;
            }
        }
    }
}

et = omp_get_wtime();
tt[bl] = et - st;
}
for (int i = 0; i < 8; i++)
{
    printf("\n\nBlock Size: %d", block[i]);
    printf("\nrun time = %f\n", tt[i]);
}
return 0;
}

```

Parallel Code:

```
#include <omp.h>
#include <stdio.h>
#include <stdlib.h>

#define MIN(X, Y) (((X) < (Y)) ? (X) : (Y))

int thread[] = {1, 2, 4, 6, 8, 10, 12, 16, 20, 32, 64, 128, 150};

int block[] = {5, 10, 50, 60, 70, 80, 90, 100};
int main()
{
    double st, et, tt[8][13];

    int omp_rank, omp_threads, L = 500, M = 500, N = 500;

    double matrix_x[500][500], matrix_y[500][500], matrix_z[500][500];

    #pragma omp for collapse(2)
    for (int i = 0; i < 500; i++)
    {
        for (int j = 0; j < 500; j++)
        {
            matrix_x[i][j] = (i * 5.123);
            matrix_y[i][j] = (j * 3.123);
            matrix_z[i][j] = 0;
        }
    }
    for(int bl = 0; bl < 4; bl++)
    {
        for (int k = 0; k < 13; k++)
        {
            omp_set_num_threads(thread[k]);
            st = omp_get_wtime();
            #pragma omp parallel private(omp_rank)
            {
                omp_rank = omp_get_thread_num();
                omp_threads = omp_get_num_threads();

                #pragma omp for
                for(int jj = 0; jj < N; jj = jj + block[bl])
                {
                    for(int kk = 0; kk < N; kk = kk + block[bl])
                    {
                        for(int i = 0; i < N; i++)
                        {
                            for(int j = jj; j < MIN(jj + block[bl], N); j++)
                            {
                                double r = 0;
                                for(int k = kk; k < MIN(kk + block[bl], N); k++)
                                {
                                    r += matrix_x[i][k] * matrix_y[k][j];
                                }
                            }
                        }
                    }
                }
            }
            et = omp_get_wtime();
            tt[bl][k] = et - st;
        }
    }
    for(int bl = 0; bl < 4; bl++)
    {
        for(int k = 0; k < 13; k++)
        {
            printf("Block %d, Thread %d, Time: %f\n", bl, k, tt[bl][k]);
        }
    }
    return 0;
}
```

```

        matrix_z[i][j] = matrix_z[i][j] + r;
    }
}
}
}
}
    et = omp_get_wtime();
    tt[bl][k] = et - st;
}
}
for (int i = 0; i < 8; i++)
{
    printf("\n\nBlock Size: %d", block[i]);
    for(int j = 0; j < 13; j++)
        printf("\nthreads = %d time = %f\n", thread[j], tt[i][j]);
}
return 0;
}

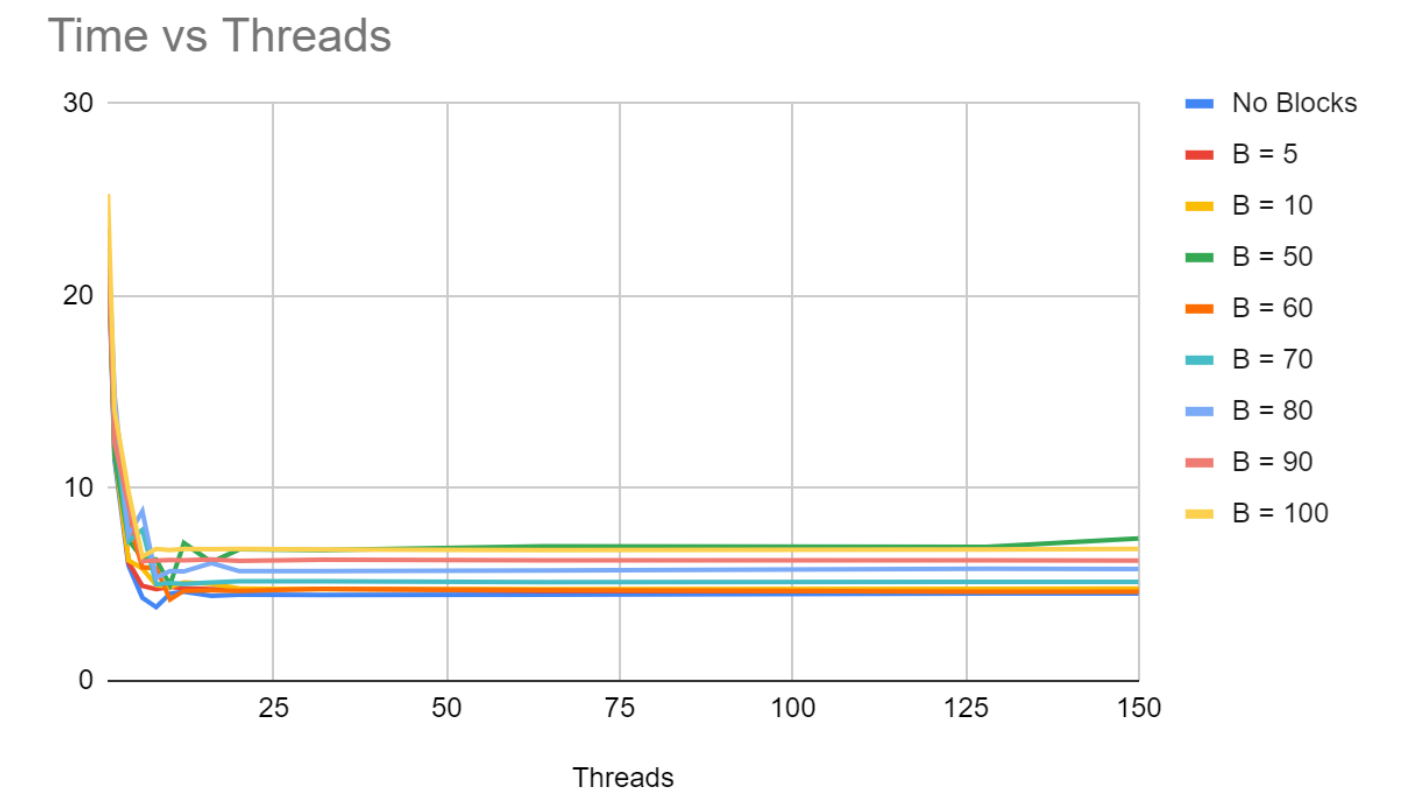
```

Analysis

- **Runtimes:**

Threads	No Blocks	B = 5	B = 10	B = 50	B = 60	B = 70	B = 80	B = 90	B = 100
1	22.83438 1	23.48643 3	23.22425 1	23.00994 5	22.9530 27	23.15141 5	23.0057 22	22.9465 44	25.30947 8
2	11.48587	12.01343 5	11.856178	11.746762	13.9065 75	13.10153 9	14.8050 51	12.37657 7	14.05080 8
4	6.000427	6.157926	6.262981	7.319166	8.40726 2	7.272827	7.587125	9.161485	9.816426
6	4.318604	4.935832	5.844545	6.397539	5.89780 8	7.838773	8.810107	6.22314	6.46305 3
8	3.824371	4.772059	4.971684	6.303351	5.85403 5	5.01490 7	5.35424 7	6.24933 8	6.853781
10	4.535858	4.881417	4.868007	4.973113	4.22913 9	5.08979	5.70299 3	6.28203 5	6.790966
12	4.62796	4.811815	5.120404	7.167268	4.68071	5.04532	5.68634	6.26236	6.851798
16	4.424744	4.766318	5.037133	6.164324	4.699177	5.11005	6.122624	6.31650 4	6.8426
20	4.478333	4.736443	4.791715	6.834658	4.68054 1	5.169086	5.701152	6.22898 1	6.855366
32	4.466217	4.787231	4.7674	6.79396	4.78444 8	5.173267	5.698317	6.29028 6	6.823517
64	4.477203	4.676044	4.76644 2	6.98979	4.718529	5.125316	5.734176	6.26024 3	6.795096

128	4.532837	4.656539	4.780197	6.967308	4.632021	5.150665	5.826881	6.265207	6.825074
150	4.538055	4.701207	4.790802	7.40836	4.634609	5.150174	5.803925	6.254129	6.85385




• Speedup:

Threads	No Blocks	B = 5	B = 10	B = 50	B = 60	B = 70	B = 80	B = 90	B = 100
1	1	1	1	1	1	1	1	1	1
2	1.988041045	1.955013949	1.958831168	1.958832996	1.650516177	1.76707599	1.553910351	1.854029915	1.801282745
4	3.805459345	3.814016765	3.708178422	3.143793296	2.730142941	3.183275912	3.03220548	2.50467517	2.578278286
6	5.287444971	4.758353404	3.973662792	3.596686945	3.891789458	2.953448837	2.611287468	3.687293553	3.916025135
8	5.970754668	4.92165604	4.671304733	3.650430541	3.92088995	4.616519309	4.296724077	3.671835961	3.692776002
10	5.034192208	4.811396568	4.770792441	4.626869528	5.427352234	4.548599255	4.033973389	3.652724635	3.726933399
12	4.934005696	4.880992515	4.535628634	3.210420623	4.903749004	4.588691104	4.045787273	3.664200717	3.69384474
16	5.160610648	4.927584144	4.610609051	3.732760478	4.884478069	4.530565259	3.757493846	3.63279181	3.698810101

20	5.098857 32	4.958664 762	4.846751 32	3.366656 386	4.903926 063	4.478821 788	4.035276 028	3.683835 928	3.691922 211
32	5.1126895 54	4.906058 011	4.871471 032	3.386823 738	4.797424 28	4.475202 034	4.037283 64	3.647933 337	3.709154 385
64	5.100144 22	5.022714 286	4.87245 0142	3.291936 525	4.864445 466	4.517070 752	4.012036 254	3.665439 824	3.724668 202
128	5.037547 346	5.04375 3096	4.858429 684	3.302558 894	4.955294 244	4.494839 987	3.948205 223	3.662535 651	3.708308 218
150	5.031755 014	4.995830 432	4.847674 982	3.105943 151	4.952527 171	4.495268 509	3.963821 379	3.669023 137	3.692738 826

Speedup vs Threads



Inference:

- Maximum speedup was observed at thread count equal to 8 in the non-blocking method.
- In the case for block based matrix multiplication: for low number of blocks (5 - 60) the maximum speedup was observed at threads counts higher than 8. For B = 80 and 90 maximum speedup was observed at thread count = 8. When B is increased further the maximum speedup was found to be at thread count = 6 (lower than 8).
- When comparing the matrix addition with and without block, lower runtime and higher speedup was observed for the method without block. Although when B = 60 the speedup was comparable at thread count = 8.